



AMAZON REDSHIFT

1.1 - 1.5

Legal Notice: This document contains privileged and/or confidential information and may not be disclosed, distributed or reproduced without the prior written permission of EPAM®.

Confidential

```
● postgresql.service - PostgreSQL database server
Loaded: loaded (/usr/lib/systemd/system/postgresql.service; enabled; preset: disabled)
Active: active (running) since Tue 2024-09-03 10:02:10 UTC; 2min 20s ago
Main PID: 2178 (postgres)
Tasks: 7 (limit: 1112)
Memory: 29.7M
CPU: 71ms
CGroup: /system.slice/postgresql.service
├─2178 /usr/bin/postgres -D /var/lib/postgresql/data
├─2199 postgres: logger "
├─2200 postgres: checkpoint "
├─2201 postgres: background writer "
├─2208 postgres: walwriter "
├─2209 postgres: autovacuum launcher "
└─2210 postgres: logical replication launcher "
```

Sep 03 10:02:10 ip-172-31-26-63.eu-central-1.compute.internal systemd[1]: Starting postgresql.service - PostgreSQL database server...

Sep 03 10:02:10 ip-172-31-26-63.eu-central-1.compute.internal postgres[2178]: 2024-09-03 10:02:10.247 UTC [2178] LOG: redirecting log output to logging collector process

Sep 03 10:02:10 ip-172-31-26-63.eu-central-1.compute.internal postgres[2178]: 2024-09-03 10:02:10.247 UTC [2178] HINT: Future log output will appear in directory "log".

Sep 03 10:02:10 ip-172-31-26-63.eu-central-1.compute.internal systemd[1]: Started postgresql.service - PostgreSQL database server.

1.6

```
[lines 1-20/20 (END)]
[ec2-user@ip-172-31-26-63 ~]$ psql -h data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-1.redshift.amazonaws.com -v
schema=public -p 5439 -U dilab_student1500 -d dev
psql: option requires an argument -- 'v'
psql: hint: try 'psql --help' for more information.
-bash: p: command not found
[ec2-user@ip-172-31-26-63 ~]$ psql -h data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-1.redshift.amazonaws.com -v schema=public -p 5439 -U dilab_student1500 -d dev
psql: warning: extra command-line argument "1.redshift.amazonaws.com" ignored
psql: error: could not translate host name "data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-" to address: Name or service not known
[ec2-user@ip-172-31-26-63 ~]$ psql -h data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-1.redshift.amazonaws.com -v schema=public -p 5439 -U dilab_student1500 -d dev
Password for user dilab_student1500:
psql (15.8, server 8.0.2)
WARNING: psql major version 15, server major version 8.0.
Some psql features might not work.
SSL connection (protocol: TLSv1.2, cipher: ECDHE-RSA-AES256-GCM-SHA384, compression: off)
Type "help" for help.

dev=>
```

1.7

```
-bash: p: Command not found
[ec2-user@ip-172-31-26-63 ~]$ psql -h data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-1.redshift.amazonaws.com -v schema=public -p 5439 -U dilab_student1500 -d dev
psql: warning: extra command-line argument "1.redshift.amazonaws.com" ignored
psql: error: could not translate host name "data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-" to address: Name or service not known
[ec2-user@ip-172-31-26-63 ~]$ psql -h data-bi-lab-redshift-cluster-3.cettexdsxw3v.eu-central-1.redshift.amazonaws.com -v schema=public -p 5439 -U dilab_student1500 -d dev
Password for user dilab_student1500:
psql (15.8, server 8.0.2)
WARNING: psql major version 15, server major version 8.0.
Some psql features might not work.
SSL connection (protocol: TLSv1.2, cipher: ECDHE-RSA-AES256-GCM-SHA384, compression: off)
Type "help" for help.

dev=> SELECT * FROM information_schema.tables LIMIT 5;
 table_catalog | table_schema | table_name | table_type | self_referencing_column_name | reference_generation | user_defined_type_catalog | user_defined_type_schema | user_defin
-----
 dev           | information_schema | information_schema_catalog_name | VIEW      |                               |                      |                          |                          |
 dev           | information_schema | applicable_roles | VIEW      |                               |                      |                          |                          |
 dev           | information_schema | check_constraints | VIEW      |                               |                      |                          |                          |
 dev           | information_schema | column_domain_usage | VIEW      |                               |                      |                          |                          |
 dev           | information_schema | column_privileges | VIEW      |                               |                      |                          |                          |
(5 rows)

dev=>
```



IGNOREHEADER 1;

```

select * from dilab_student1500.dim_stores

SELECT
  "column" AS column_name,
  "type" AS data_type,
  encoding
FROM
  pg_table_def
WHERE
  tablename = 'dim_customers_scd'
  AND schemaname = 'dilab_student1500';

```

pg_table_def 1 X

SELECT "column" AS column_name, "type" AS data_type, encoding

Grid

	A-Z column_name	data_type	A-Z encoding
1	customer_surr_id	bigint	az64
2	f_name	character varying(50)	lzo
3	l_name	character varying(50)	lzo
4	email	character varying(100)	lzo
5	cust_phone	character varying(30)	lzo
6	start_dt	date	az64
7	end_dt	date	az64
8	is_active	character varying(10)	lzo
9	insert_dt	date	az64
10	source_id	character varying(100)	lzo
11	source_system	character varying(100)	lzo
12	source_entity	character varying(100)	lzo

Record

Refresh Save Cancel

Export data

AMT en_US Writable Smart Insert 58:1 [189] Sel

ANALYZE COMPRESSION dilab_student1500.dim_customers_scd_withoutcomp;

Results 1 x

ANALYZE COMPRESSION dilab_student1500

A-z Table	A-z Column	A-z Encoding	A-z Est_reduction_pct
1 dim_customers_scd_withoutcomp	customer_surr_id	raw	0.00
2 dim_customers_scd_withoutcomp	f_name	raw	0.00
3 dim_customers_scd_withoutcomp	l_name	raw	0.00
4 dim_customers_scd_withoutcomp	email	raw	0.00
5 dim_customers_scd_withoutcomp	cust_phone	raw	0.00
6 dim_customers_scd_withoutcomp	start_dt	raw	0.00
7 dim_customers_scd_withoutcomp	end_dt	raw	0.00
8 dim_customers_scd_withoutcomp	is_active	raw	0.00
9 dim_customers_scd_withoutcomp	insert_dt	raw	0.00
10 dim_customers_scd_withoutcomp	source_id	raw	0.00
11 dim_customers_scd_withoutcomp	source_system	raw	0.00
12 dim_customers_scd_withoutcomp	source_entity	raw	0.00

```
analyze compression dilab_student1500.dim_customers_scd_analyzedcomp;

-- Insert data from the original table into the newly created table
INSERT INTO dilab_student1500.dim_customers_scd_analyzedcomp
SELECT * FROM dilab_student1500.dim_customers_scd_withoutcomp;

CREATE TABLE dilab_student1500.dim_customers_scd_analyzedcomp (
  customer_surr_id INT ENCODE AZ64, -- Replace with recommended encoding
  f_name VARCHAR(50) ENCODE ZSTD, -- Replace with recommended encoding
  l_name VARCHAR(50) ENCODE ZSTD, -- Replace with recommended encoding
  email VARCHAR(100) ENCODE LZ0, -- Replace with recommended encoding
  cust_phone VARCHAR(20) ENCODE ZSTD, -- Replace with recommended encoding
  start_dt DATE ENCODE ZSTD, -- Replace with recommended encoding
  end_dt DATE ENCODE ZSTD, -- Replace with recommended encoding
  is_active BOOLEAN ENCODE RAW, -- Replace with recommended encoding
  insert_dt TIMESTAMP ENCODE ZSTD, -- Replace with recommended encoding
  source_id INT ENCODE AZ64, -- Replace with recommended encoding
  source_system VARCHAR(50) ENCODE ZSTD -- Replace with recommended encoding
);
```

Results 1 x Statistics 0

analyze compression dilab_student1500.dim

A-z Table	A-z Column	A-z Encoding	A-z Est_reduction_pct
1 dim_customers_scd_analyzedcomp	customer_surr_id	raw	0.00
2 dim_customers_scd_analyzedcomp	f_name	zstd	0.00
3 dim_customers_scd_analyzedcomp	l_name	raw	0.00
4 dim_customers_scd_analyzedcomp	email	raw	0.00
5 dim_customers_scd_analyzedcomp	cust_phone	raw	0.00
6 dim_customers_scd_analyzedcomp	start_dt	raw	0.00
7 dim_customers_scd_analyzedcomp	end_dt	raw	0.00
8 dim_customers_scd_analyzedcomp	is_active	raw	0.00
9 dim_customers_scd_analyzedcomp	insert_dt	raw	0.00
10 dim_customers_scd_analyzedcomp	source_id	raw	0.00
11 dim_customers_scd_analyzedcomp	source_system	raw	0.00
12 dim_customers_scd_analyzedcomp	source_entity	raw	0.00

[illegible]