



PG_DWH TASK 3

Legal Notice: This document contains privileged and/or confidential information and may not be disclosed, distributed or reproduced without the prior written permission of EPAM®.

Confidential

```
Query  Query History
1  DROP TABLE IF EXISTS labs.person;
2
3  SET search_path TO labs;
4
5  CREATE TABLE labs.person (
6  id integer NOT NULL,
7  name varchar(15)
8  );
9  INSERT INTO person VALUES(1, 'Bob');
10 INSERT INTO person VALUES(2, 'Alice');
11 INSERT INTO person VALUES(3, 'Robert');
12
13
14

Data Output  Messages  Notifications
NOTICE:  table "person" does not exist, skipping
INSERT 0 1

Query returned successfully in 214 msec.
```

Table created , values are inserted.

```
3 SET search_path TO labs;
4
5 CREATE TABLE labs.person (
6   id integer NOT NULL,
7   name varchar(15)
8 );
9 INSERT INTO person VALUES(1, 'Bob');
10 INSERT INTO person VALUES(2, 'Alice');
11 INSERT INTO person VALUES(3, 'Robert');
12
13
14 CREATE TABLE IF NOT EXISTS labs.test_simple(a int,b int);
15
16 insert into test_simple values (generate_series(1,1000000));
```

Data Output Messages Notifications

INSERT 0 1000000

Query returned successfully in 2 secs 353 msec.

Total rows: 3 of 3 Query complete 00:00:02.353

Simple table created and populated in 2 seconds in first query.

```
5 CREATE TABLE labs.person (
6   id integer NOT NULL,
7   name varchar(15)
8 );
9 INSERT INTO person VALUES(1, 'Bob');
10 INSERT INTO person VALUES(2, 'Alice');
11 INSERT INTO person VALUES(3, 'Robert');
12
13
14 CREATE TABLE IF NOT EXISTS labs.test_simple(a int,b int);
15
16 insert into test_simple values (generate_series(1,1000000));
17
18
19 insert into test_simple values (generate_series(1,5000000));
```

Data Output Messages Notifications

INSERT 0 5000000

Query returned successfully in 11 secs 619 msec.

Total rows: 3 of 3 Query complete 00:00:11.619

Second query took 11 seconds to execute.

```
9  --INSERT INTO person VALUES(1, 'Bob');
10 --INSERT INTO person VALUES(2, 'Alice');
11 --INSERT INTO person VALUES(3, 'Robert');
12
13
14 CREATE TABLE IF NOT EXISTS labs.test_simple(a int,b int);
15
16 --insert into test_simple values (generate_series(1,1000000));
17 --insert into test_simple values (generate_series(1,5000000));
18
19
20 CREATE UNLOGGED TABLE IF NOT EXISTS labs.test_unlogged(a int, b int);
21 insert into test_unlogged values (generate_series(1,1000000));
22
23
```

Data Output Messages Notifications

INSERT 0 1000000

Query returned successfully in 1 secs 316 msec.

Total rows: 3 of 3 Query complete 00:00:01.316

Unlogged table took 1.3 second to be populated 1.000.000 rows.

```
9  --INSERT INTO person VALUES(1, 'Bob');
10 --INSERT INTO person VALUES(2, 'Alice');
11 --INSERT INTO person VALUES(3, 'Robert');
12
13
14 CREATE TABLE IF NOT EXISTS labs.test_simple(a int,b int);
15
16 --insert into test_simple values (generate_series(1,1000000));
17 --insert into test_simple values (generate_series(1,5000000));
18
19
20 CREATE UNLOGGED TABLE IF NOT EXISTS labs.test_unlogged(a int, b int);
21 insert into test_unlogged values (generate_series(1,1000000));
22 insert into test_simple values (generate_series(1,5000000));
23
```

Data Output Messages Notifications

INSERT 0 5000000

Query returned successfully in 11 secs 114 msec.

Total rows: 3 of 3 Query complete 00:00:11.114

And second one took 11 second , as we see unlogged tables are faster than usual tables.

```
23
24
25
26 CREATE TABLE labs.users(user_id INT GENERATED BY DEFAULT AS IDENTITY
27 PRIMARY KEY, login VARCHAR(30)) INHERITS (person);
28
29
```

Data Output Messages Notifications

CREATE TABLE

Query returned successfully in 154 msec.

Table created .

```
23
24
25
26 CREATE TABLE labs.users(user_id INT GENERATED BY DEFAULT AS IDENTITY
27 PRIMARY KEY, login VARCHAR(30)) INHERITS (person);
28
29 INSERT INTO users VALUES (1, 'new Bob', 'NewLogin');
30 INSERT INTO users VALUES (999, 'TestUser', 'TestLogin');
```

Data Output Messages Notifications

ERROR: invalid input syntax for type integer: "NewLogin"
LINE 2: INSERT INTO users VALUES (1, 'new Bob', 'NewLogin');
A

SQL state: 22P02
Character: 42

Third value need to be an integer as its inherits from person.

```
25
26 CREATE TABLE labs.users(user_id INT GENERATED BY DEFAULT AS IDENTITY
27 PRIMARY KEY, login VARCHAR(30)) INHERITS (person);
28
29 INSERT INTO users VALUES (5, 'new Bob', 123, 'logged_in');
30 INSERT INTO users VALUES (999, 'TestUser', 'TestLogin');
```

Data Output Messages Notifications

INSERT 0 1

Query returned successfully in 123 msec.

Inserted new row , so 5 and 'new Bob' values are going to be values for person table also ,
so it's insering there as well.

Data Output					Messages		Notifications	
	Id integer		name character varying (15)	user_Id [PK] integer		login character varying (30)		
1		1	new Bob		1	[null]		
2		5	new Bob		123	logged_in		
3		1	new Bob		454	[null]		
4		4	new Bob		666	[null]		












Users table after inserting and person table .












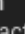

	Id integer		name character varying (15)
1		1	Bob
2		2	Alice
3		3	Robert
4		1	new Bob
5		1	new Bob
6		4	new Bob
7		5	new Bob

In other words users table will contain person table in it.

31	
32	
33	
34	UPDATE person
35	SET name = 'not Bob'
36	where id = 1;
Data Output	
Messages	
UPDATE 3	
Query returned successfully in 237 msec.	

Update statement updates all new Bobs to not Bobs in both tables .

Data Output Messages Notifications			
	        		
	id integer		name character varying (15) 
1		2	Alice
2		3	Robert
3		1	not Bob
4		4	new Bob
5		5	new Bob
6		1	not Bob
7		1	not Bob

Data Output Messages Notifications				
	        			
	id integer 	name character varying (15) 	user_id [PK] integer 	login character varying (30) 
1	1	not Bob	1	[null]
2	5	new Bob	123	logged_in
3	1	not Bob	454	[null]
4	4	new Bob	666	[null]

Alter person.

37	
38	
39	ALTER TABLE person ADD COLUMN status integer DEFAULT 0;
40	ALTER TABLE person ADD CONSTRAINT status CHECK (status in (0,1)) NO
41	INHERIT;
42	ALTER TABLE person ADD CONSTRAINT id UNIQUE (id, name);
Data Output Messages Notifications	
ALTER TABLE	
Query returned successfully in 175 msec.	

We can see that by Alterin person table users table were altered as well

Data Output Messages Notifications						
	Id integer	name character varying (15)	user_id [PK] integer	login character varying (30)	status integer	
1	1	not Bob	1	[null]	0	
2	5	new Bob	123	logged_in	0	
3	1	not Bob	454	[null]	0	
4	4	new Bob	666	[null]	0	

- 2.
- Table created and populated with 15776640 values.
666 MB's

51	
52	
53	INSERT INTO test_index(num, load_date)
54	SELECT random(), x
55	FROM generate_series('2017-01-01 0:00'::timestampz,
56	'2021-12-31 23:59:59'::timestampz, '10 seconds'::interval) x;
57	
58	
59	SELECT pg_size_pretty(pg_relation_size('test_index'));
Data Output Messages Notifications	
	pg_size_pretty text
1	666 MB


```
60
61
62
63 SELECT date_trunc('year', load_date), max(num)
64 FROM test_index
65 WHERE load_date BETWEEN '2021-09-01 0:00' AND '2021-10-31 11:59:59'
66 GROUP BY 1
67 ORDER BY 1;
```

Data Output Messages Notifications

	date_trunc timestamp with time zone	max double precision
1	2021-01-01 00:00:00+04	0.9999981809137353

Total rows: 1 of 1 Query complete 00:00:02.612

Took 2.6seconds.

```
67 ORDER BY 1;
68
69
70 CREATE INDEX test_index_date ON test_index (load_date);
71
```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 17 secs 484 msec.

Creating index took 17.4 seconds and size is 338MB.

```
68
69
70 CREATE INDEX test_index_date ON test_index (load_date);
71 SELECT pg_size_pretty(pg_relation_size('test_index_date'));
72
```

Data Output Messages Notifications

	pg_size_pretty text
1	338 MB

```
71 SELECT pg_size_pretty(pg_relation_size('test_index_date'));
72 DROP INDEX test_index_date;
73 CREATE INDEX test_index_date ON test_index (load_date);
74 CREATE INDEX test_index_load_date_brin ON test_index USING BRIN (load_date);
75
```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 7 secs 682 msec.

BRIN index created took 7.6 sec and it's just 40kb.

```
73 CREATE INDEX test_index_date ON test_index (load_date);
74 CREATE INDEX test_index_load_date_brin ON test_index USING BRIN (load_date);
75 SELECT pg_size_pretty(pg_relation_size('test_index_load_date_brin'));
76
```

Data Output Messages Notifications

	pg_size_pretty text
1	40 kB

3.

Table created and populated with 10.000.000 values , size 651 MB

```
81
82 CREATE TABLE test_index AS SELECT id, md5(id::text) as t_hash
83 FROM generate_series(1, 10000000) AS id;
84
85 SELECT pg_size_pretty(pg_relation_size('test_index'));
86
```

Data Output Messages Notifications

	pg_size_pretty text
1	651 MB

```
78 DROP TABLE test_index;
79
80
81
82 CREATE TABLE test_index AS SELECT id, md5(id::text) as t_hash
83 FROM generate_series(1, 10000000) AS id;
84
85 SELECT pg_size_pretty(pg_relation_size('test_index'));
86 SELECT * FROM test_index WHERE t_hash LIKE '%ceea167a5a%';
```

Data Output Messages Notifications

	id Integer	t_hash text
1	7	8f14e45fcea167a5a36dedd4bea2543

Total rows: 1 of 1 Query complete 00:00:02.509

2.5 second to find pattern.

```
82 CREATE TABLE test_index AS SELECT id, md5(id::text) as t_hash
83 FROM generate_series(1, 10000000) AS id;
84
85 SELECT pg_size_pretty(pg_relation_size('test_index'));
86 SELECT * FROM test_index WHERE t_hash LIKE '%ceea167a5a%';
87
88 CREATE EXTENSION pg_trgm;
89 CREATE INDEX idx_text_index_gist ON test_index USING gist(t_hash
90 gist_trgm_ops);
```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 13 min 35 secs.

Total rows: 1 of 1 Query complete 00:13:35.696

Creation gist index took 13 minutes :(

```
84
85 SELECT pg_size_pretty(pg_relation_size('test_index'));
86 SELECT * FROM test_index WHERE t_hash LIKE '%ceea167a5a%';
87
88 CREATE EXTENSION pg_trgm;
89 CREATE INDEX idx_text_index_gist ON test_index USING gist(t_hash
90 gist_trgm_ops);
91
92
93 DROP INDEX idx_text_index_gist;
94
95
96 CREATE INDEX idx_text_index_gin ON test_index USING gin (t_hash
97 gin_trgm_ops);
```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 4 min 31 secs.

Total rows: 1 of 1 Query complete 00:04:31.598

With gin index it took 4minutes and size is 613 MB

```
84
85 SELECT pg_size_pretty(pg_relation_size('test_index'));
86 SELECT * FROM test_index WHERE t_hash LIKE '%ceea167a5a%';
87
88 CREATE EXTENSION pg_trgm;
89 CREATE INDEX idx_text_index_gist ON test_index USING gist(t_hash
90 gist_trgm_ops);
91
92
93 DROP INDEX idx_text_index_gist;
94
95
96 CREATE INDEX idx_text_index_gin ON test_index USING gin (t_hash
97 gin_trgm_ops);
```

Data Output Messages Notifications

CREATE INDEX

Query returned successfully in 4 min 31 secs.

Total rows: 1 of 1 Query complete 00:04:31.598

```
91  
92  
93 DROP INDEX idx_text_index_gist;  
94  
95  
96 CREATE INDEX idx_text_index_gin ON test_index USING gin (t_hash  
97 gin_trgm_ops);  
98  
99 SELECT pg_size_pretty(pg_relation_size('idx_text_index_gin'));  
100 DROP TABLE test_index  
101  
102 |  
103 CREATE TABLE test_index AS SELECT id, md5(id::text) as t_hash  
104 FROM generate_series(1, 10000000) AS id;
```

Data Output Messages Notifications

SELECT 10000000

Query returned successfully in 1 min 4 secs.

Total rows: 10 of 10 Query complete 00:01:04.318

Population table this time took 1 minute which is way faster then first time.

4.

```
113  
114  
115 CREATE FOREIGN TABLE labs.test_foreign_table  
116 (  
117 LatD INT,  
118 LatM INT,  
119 LatS INT,  
120 NS TEXT,  
121 LonD INT,  
122 LonM INT,  
123 LonS INT,  
124 EW TEXT,  
125 City TEXT,  
126 State TEXT  
127 )  
128 SERVER test_import  
129 OPTIONS  
130 (  
131 filename '/home/styop/Downloads/ities_list.csv',  
132 format 'csv',  
133 header 'true',  
134 delimiter ',',  
135 );
```

Data Output Messages Notifications

CREATE FOREIGN TABLE

Query returned successfully in 115 msec.

Foreign table created using csv file.

```

125 City TEXT,
126 State TEXT
127 )
128 SERVER test_import
129 OPTIONS
130 (
131   filename '/home/styop/Downloads/cities_list.csv',
132   format 'csv',
133   header 'true',
134   delimiter ','
135 );
136
137
138 SELECT * FROM test_foreign_table;

```

Data Output Messages Notifications

	latd integer	latm integer	lats integer	ns text	lond integer	lonm integer	lons integer	ew text	city text	state text
1	41	5	59	N	80	39	0	W	Youngstown	OH
2	42	52	48	N	97	23	23	W	Yankton	SD
3	46	35	59	N	120	30	36	W	Yakima	WA
4	42	16	12	N	71	48	0	W	Worcester	MA
5	43	37	48	N	89	46	11	W	WisconsinDells	WI
6	36	5	59	N	80	15	0	W	Winston-Salem	NC
7	49	52	48	N	97	9	0	W	Winnipeg	MB
8	39	11	23	N	78	9	36	W	Winchester	VA
9	34	14	24	N	77	55	11	W	Wilmington	NC
10	39	45	0	N	75	33	0	W	Wilmington	DE
11	48	9	0	N	103	37	12	W	Williston	ND

Total rows: 128 of 128 Query complete 00:00:00.134

There is 128 rows.

```

130 (
131   filename '/home/styop/Downloads/cities_list.csv',
132   format 'csv',
133   header 'true',
134   delimiter ','
135 );
136
137
138 SELECT * FROM test_foreign_table;
139
140 SELECT count(*) FROM test_foreign_table;
141
142
143 CREATE MATERIALIZED VIEW mview as ( SELECT * FROM test_foreign_table);

```

Data Output Messages Notifications

SELECT 128

Query returned successfully in 260 msec.

Total rows: 1 of 1 Query complete 00:00:00.260

Materialized view created with 128 rows.

I deleted 2 rows from file but the count in materialized view remains same, as we store our previous data in materialized view (it means we copy the file in the view)

```
132 format 'csv',
133 header 'true',
134 delimiter ','
135 );
136
137
138 SELECT * FROM test_foreign_table;
139
140 SELECT count(*) FROM test_foreign_table;
141
142
143 CREATE MATERIALIZED VIEW mview as ( SELECT * FROM test_foreign_table);
144 select count(*) from mview
145
```

Data Output Messages Notifications

	count bigint	
1	128	

Total rows: 1 of 1 Query complete 00:00:00.210