



PG_DWH TASK 5

Legal Notice: This document contains privileged and/or confidential information and may not be disclosed, distributed or reproduced without the prior written permission of EPAM®.

Confidential

QueryQuery History

```
1 CREATE TABLE test_joins_a
2 (
3   id1 int,
4   id2 int
5 );
6 CREATE TABLE test_joins_b
7 (
8   id1 int,
9   id2 int
10 );
11 INSERT INTO test_joins_a values(generate_series(1,10000),3);
12 INSERT INTO test_joins_b values(generate_series(1,10000),3);
13 ANALYZE;
14
15
16 SELECT * FROM test_joins_a a, test_joins_b b
17 WHERE a.id1 > b.id1;
```

Data OutputMessagesNotifications

	id1 integer		id2 integer		id1 integer		id2 integer	
1		2		3		1		3
2		3		3		1		3
3		3		3		2		3
4		4		3		1		3
5		4		3		2		3
6		4		3		3		3
7		5		3		1		3
8		5		3		2		3
9		5		3		3		3
Total rows: 1000 of 49995000				Query complete 00:01:28.483				

Table created and join query executed.

17 **WHERE** a.id1 > b.id1;

Data Output Messages Notifications

QUERY PLAN text

1	Nested Loop (cost=0.00..1500315.00 rows=33333333 width=16) (actual time=130.571..25322.838 rows=49995000 loop...
2	Join Filter: (a.id1 > b.id1)
3	Rows Removed by Join Filter: 50005000
4	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=8) (actual time=0.020..17.406 rows=10000 loops=1)
5	-> Materialize (cost=0.00..195.00 rows=10000 width=8) (actual time=0.000..1.069 rows=10000 loops=10000)
6	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=8) (actual time=0.007..1.466 rows=10000 loop...
7	Planning Time: 0.107 ms
8	JIT:
9	Functions: 6
10	Options: Inlining true, Optimization true, Expressions true, Deforming true
11	Timing: Generation 0.944 ms, Inlining 4.177 ms, Optimization 85.505 ms, Emission 35.745 ms, Total 126.371 ms
12	Execution Time: 28493.942 ms

Total rows: 12 of 12 Query complete 00:00:28.747

As we see postgres choose to Sequential Scan method for this join query , it is firstly scanned thru test_joins_a and then test_joins_b. For each row in test_joins_a , Iterate through the materialized rows from test_joins_b (10,000 rows for each row in test_joins_a). Apply the join filter a.id1 > b.id1 to each pair of rows.

19
20 **EXPLAIN ANALYZE**
21
22 **SELECT** *
23 **FROM** test_joins_a a
24 **CROSS JOIN** test_joins_b b;

Data Output Messages Notifications

QUERY PLAN text

1	Nested Loop (cost=0.00..1250315.00 rows=100000000 width=16) (actual time=77.734..31379.757 rows=100000000 loo...
2	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=8) (actual time=0.016..21.627 rows=10000 loops=1)
3	-> Materialize (cost=0.00..195.00 rows=10000 width=8) (actual time=0.000..1.217 rows=10000 loops=10000)
4	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=8) (actual time=0.005..1.327 rows=10000 loop...
5	Planning Time: 0.115 ms
6	JIT:
7	Functions: 3
8	Options: Inlining true, Optimization true, Expressions true, Deforming true
9	Timing: Generation 0.401 ms, Inlining 2.930 ms, Optimization 51.328 ms, Emission 23.438 ms, Total 78.097 ms
10	Execution Time: 38241.348 ms

As we can see Cross join exequates exactly in the same method and order.

QueryQuery History

```
28 SET enable_mergejoin = off;
29
30 EXPLAIN ANALYZE
31 SELECT *
32 FROM test_joins_a a
33 JOIN test_joins_b b
34 ON a.id1 > b.id1;
35
36
37 EXPLAIN ANALYZE
38 SELECT *
39 FROM test_joins_a a
40 WHERE EXISTS (
41     SELECT 1
42     FROM test_joins_b b
43     WHERE a.id1 > b.id1
44 );
45
```

Data OutputMessagesNotifications

QUERY PLAN

text

1	Nested Loop Semi Join (cost=10000000000.00..10001000464.98 rows=3333 width=8) (actual time=130.798..137.371 rows=9...
2	Join Filter: (a.id1 > b.id1)
3	Rows Removed by Join Filter: 10000
4	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=8) (actual time=0.021..1.727 rows=10000 loops=1)
5	-> Materialize (cost=0.00..195.00 rows=10000 width=4) (actual time=0.012..0.013 rows=2 loops=10000)
6	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=4) (actual time=123.029..125.125 rows=10000 loop...
7	Planning Time: 0.121 ms
8	JIT:
9	Functions: 7
10	Options: Inlining true, Optimization true, Expressions true, Deforming true
11	Timing: Generation 0.913 ms, Inlining 6.128 ms, Optimization 76.978 ms, Emission 39.946 ms, Total 123.965 ms
12	Execution Time: 139.183 ms

Semi Join query created.

29

```
30
31 SET enable_hashjoin = on;
32
33
34
35 EXPLAIN ANALYZE
36 SELECT *
37 FROM test_joins_a a
38 WHERE EXISTS (
39     SELECT 1
40     FROM test_joins_b b
41     WHERE a.id1 > b.id1
42 );
43
```

Data OutputMessagesNotifications

QUERY PLAN

text

1	Nested Loop Semi Join (cost=0.00..1000464.98 rows=3333 width=8) (actual time=144.244..149.391 rows=9999 loops=1)
2	Join Filter: (a.id1 > b.id1)
3	Rows Removed by Join Filter: 10000
4	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=8) (actual time=0.020..1.348 rows=10000 loops=1)
5	-> Materialize (cost=0.00..195.00 rows=10000 width=4) (actual time=0.014..0.014 rows=2 loops=10000)
6	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=4) (actual time=138.540..140.220 rows=10000 loop...
7	Planning Time: 0.126 ms
8	JIT:
9	Functions: 7
10	Options: Inlining true, Optimization true, Expressions true, Deforming true
11	Timing: Generation 0.978 ms, Inlining 6.083 ms, Optimization 95.499 ms, Emission 36.978 ms, Total 139.539 ms
12	Execution Time: 151.196 ms

enable_hashjoin set to off and query plan executed , then set to on.

Query

Query History

```
40 FROM test_joins_b b
41 WHERE a.id1 > b.id1
42 );
43 SET enable_hashjoin = off;
44
45
46 -- Ensure columns are indexed or sorted
47 CREATE INDEX idx_a_id1 ON test_joins_a(id1);
48 CREATE INDEX idx_b_id1 ON test_joins_b(id1);
49
50 -- Query to force MERGE JOIN
51 EXPLAIN ANALYZE
52 SELECT *
53 FROM test_joins_a a
54 JOIN test_joins_b b
55 ON a.id1 = b.id1;
56
57
```

Data Output

Messages

Notifications

+

📄

▼

📄

▼

🗑️

📄

📄

📄

QUERY PLAN

text

1	Merge Join (cost=0.57..786.57 rows=10000 width=16) (actual time=0.047..21.187 rows=10000 loops=1)
2	Merge Cond: (a.id1 = b.id1)
3	-> Index Scan using idx_a_id1 on test_joins_a a (cost=0.29..318.29 rows=10000 width=8) (actual time=0.023..5.227 rows=10000 loop...)
4	-> Index Scan using idx_b_id1 on test_joins_b b (cost=0.29..318.29 rows=10000 width=8) (actual time=0.015..6.123 rows=10000 loop...)
5	Planning Time: 0.373 ms
6	Execution Time: 22.423 ms

Query for Merge join , first I disabled hashjoin then created index on both tables(make sure the join condition columns are sorted) and execute query.

```
42 CREATE INDEX idx_a_id1 ON test_joins_a(id1);
43 CREATE INDEX idx_b_id1 ON test_joins_b(id1);
44
45
46 SET enable_mergejoin = off;
47 SET enable_nestloop = off;
48 SET enable_mergejoin = on;
49 SET enable_hashjoin = on;
50
51
52
53 EXPLAIN ANALYZE
54 SELECT *
55 FROM test_joins_a a
56 JOIN test_joins_b b
57 ON a.id1 = b.id1;
58
59
```

Data Output Messages Notifications

QUERY PLAN

	text
1	Hash Join (cost=270.00..552.50 rows=10000 width=16) (actual time=4.770..13.891 rows=10000 loops=1)
2	Hash Cond: (a.id1 = b.id1)
3	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=8) (actual time=0.016..1.862 rows=10000 loops=1)
4	-> Hash (cost=145.00..145.00 rows=10000 width=8) (actual time=4.728..4.730 rows=10000 loops=1)
5	Buckets: 16384 Batches: 1 Memory Usage: 519kB
6	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=8) (actual time=0.007..1.592 rows=10000 loop...
7	Planning Time: 0.314 ms
8	Execution Time: 14.825 ms

After disabling merge join we can see that postgres choose hash join method for execution.

```
58
59
60
61 CREATE TABLE test_joins_c
62 (
63 id1 int,
64 id2 int
65 );
66 INSERT INTO test_joins_c
67 values(generate_series(1,1000000),(random()*10)::int);
68
69
70
71 EXPLAIN
72 SELECT c.id2
73 FROM test_joins_b b
74 JOIN test_joins_a a on (b.id1 = a.id1)
75 LEFT JOIN test_joins_c c on (c.id1 = b.id1);
```

Data Output Messages Notifications

QUERY PLAN

	text
1	Hash Right Join (cost=677.50..28853.69 rows=1000050 width=4)
2	Hash Cond: (c.id1 = b.id1)
3	-> Seq Scan on test_joins_c c (cost=0.00..14425.50 rows=1000050 width=8)
4	-> Hash (cost=552.50..552.50 rows=10000 width=4)
5	-> Hash Join (cost=270.00..552.50 rows=10000 width=4)
6	Hash Cond: (b.id1 = a.id1)
7	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=4)
8	-> Hash (cost=145.00..145.00 rows=10000 width=4)
9	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=4)

Table created and in query plan we see that first it sequentially scan on test_joins_a table , then Hashing it , next it goes to other table sequentially then hashed it too , then hash join between tables.After than it starts same process with third table , first sequentially then hash join with other joined table.


```
60
61 CREATE TABLE test_joins_c
62 (
63   id1 int,
64   id2 int
65 );
66 INSERT INTO test_joins_c
67 values(generate_series(1,1000000),(random()*10)::int);
68
69 SET join_collapse_limit = 1
70
71 EXPLAIN
72 SELECT c.id2
73 FROM test_joins_b b
74 JOIN test_joins_a a on (b.id1 = a.id1)
75 LEFT JOIN test_joins_c c on (c.id1 = b.id1);
```

Data Output Messages Notifications

QUERY PLAN

	text
1	Hash Right Join (cost=677.50..18952.50 rows=10000 width=4)
2	Hash Cond: (c.id1 = b.id1)
3	-> Seq Scan on test_joins_c c (cost=0.00..14425.00 rows=1000000 width=8)
4	-> Hash (cost=552.50..552.50 rows=10000 width=4)
5	-> Hash Join (cost=270.00..552.50 rows=10000 width=4)
6	Hash Cond: (b.id1 = a.id1)
7	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=4)
8	-> Hash (cost=145.00..145.00 rows=10000 width=4)
9	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=4)

join_collapse_limit set to 1 , we see that it start to use Hash Right join for last step.


```
67 values(generate_series(1,100000),(Random()*10)::int),
68
69 SET join_collapse_limit = 8;
70
71 EXPLAIN
72 SELECT c.id2
73 FROM test_joins_b b
74 JOIN test_joins_a a on (b.id1 = a.id1)
75 LEFT JOIN test_joins_c c on (c.id1 = b.id1);
```

Data Output Messages Notifications

QUERY PLAN

	text
1	Hash Join (cost=540.00..18952.50 rows=10000 width=4)
2	Hash Cond: (b.id1 = a.id1)
3	-> Hash Right Join (cost=270.00..18545.00 rows=10000 width=8)
4	Hash Cond: (c.id1 = b.id1)
5	-> Seq Scan on test_joins_c c (cost=0.00..14425.00 rows=1000000 width=4)
6	-> Hash (cost=145.00..145.00 rows=10000 width=4)
7	-> Seq Scan on test_joins_b b (cost=0.00..145.00 rows=10000 width=4)
8	-> Hash (cost=145.00..145.00 rows=10000 width=4)
9	-> Seq Scan on test_joins_a a (cost=0.00..145.00 rows=10000 width=4)

Returned to 8.

```
97
98
99
100 SELECT s.store_id,s.store_name,o.order_id,o.order_cost,o.order_num
101 FROM stores s
102 CROSS JOIN LATERAL (
103     SELECT order_id,order_cost,order_num FROM orders o
104     WHERE o.order_cost < s.max_order_cost
105     ORDER BY o.order_cost DESC
106     LIMIT 10
107 ) o;
108
```

Data Output Messages Notifications

	store_id integer	store_name text	order_id integer	order_cost integer	order_num text
1	1	grossery shop	259	754	order number 259
2	1	grossery shop	10	745	order number 10
3	1	grossery shop	28	744	order number 28
4	1	grossery shop	24	723	order number 24
5	1	grossery shop	180	702	order number 180
6	1	grossery shop	559	667	order number 559
7	1	grossery shop	26	656	order number 26
8	1	grossery shop	275	635	order number 275
9	1	grossery shop	30	617	order number 30
10	1	grossery shop	114	612	order number 114
11	2	bakerv	12	94	order number 12

Total rows: 25 of 25 Query complete 00:00:00.240

S

Query for TOP 10 of orders by cost for each store used LATERAL JOIN.

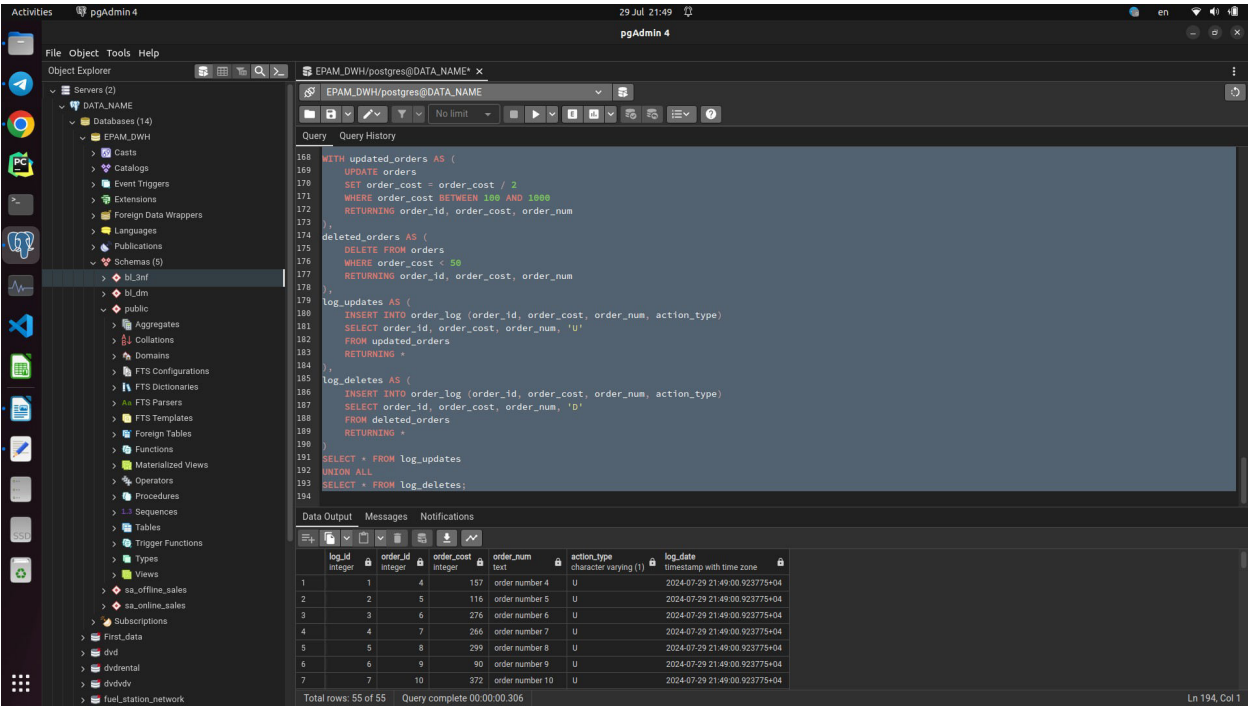
QueryQuery History

```
130
131 WITH RECURSIVE test_recursive AS (
132     SELECT e.emp_id,e.emp_name,e.manager_id,NULL::VARCHAR AS manager_name,1 AS level
133     FROM emp e
134     WHERE e.manager_id IS NULL
135
136     UNION ALL
137     SELECT e.emp_id,e.emp_name,e.manager_id,eh.emp_name AS manager_name,eh.level + 1 AS level
138     FROM emp e
139     JOIN test_recursive eh ON e.manager_id = eh.emp_id
140 )
141
142 -- Select final result
143 SELECT
144     emp_id,
145     emp_name,
146     manager_name,
147     level
148 FROM
149     test_recursive
150 ORDER BY
151     level, emp_id;
152
```

Data OutputMessagesNotifications

	emp_id integer	emp_name character varying (100)	manager_name character varying	level integer
1	1	Alice	[null]	1
2	2	Bob	Alice	2
3	5	Eve	Alice	2
4	3	Charlie	Bob	3
5	4	David	Bob	3
6	6	Frank	Eve	3
7	7	Grace	Eve	3

Recursive query for selectiong employees.



Changing data using CTE's ,